

HEMMINGWAY-1

Competent prose. Uneven follow-through.

NO COMMUNITY RANK YET / MODEL ID: hemmingway_1

01 / RESPONSE COVERAGE

100%

220 / 220 turns answered

Full response coverage across 20 sessions.

02 / FAILURE MODES

2.3%

Agency violations

44 probes / interval 0.4-11.8%

0.0%

POV / tense breaks

33 probes / interval 0.0-10.4%

NINE OTHER TRAP MODES

6 / 34

probes failed (17.6%)

Interval 8.3-33.5%

03 / WILLINGNESS & JUDGMENT

ROUND 04

-0.35

Youden's J

Over-refusal includes soft deflection, not just explicit refusals.
Boundaries: non-consent + frame-break only. Not a general safety rating.

25% held the tested boundary

60% over-refusal on harder requests

Intimacy 67% / Gore 10%

04 / BEHAVIOR

296 words per reply

Roster average: 320 words

Phrase repetition 0.057

Roster average: 0.060

05 / CRAFT

Recycled descriptions.

Narrated emotions.

Agency violations.

Top flaws / single-rater v2 / 20 sessions

06 / PRODUCTION DEFECTS

0% leaks & degenerate repetition

1.8% wrote the user's turn (4/220)

1.0x token overhead

Billed per visible character / relative floor

07 / SUBJECTIVE

Judge-sensitive, not a rank.

Single-judge estimates; other judges can reorder close models.

No human arena data for this model.

THE TAKEAWAY

Competent prose, but the session-level judge flags recurring descriptions and scene beats, alongside instruction and detail-retention slips. Frequent intimacy deflection coexists with weak adherence to the tested hard boundaries.